

# Yue Zhao

---

## CONTACT INFORMATION

✉ [yue.z@usc.edu](mailto:yue.z@usc.edu)  
🐙 [github.com/yzhao062](https://github.com/yzhao062)  
in [linkedin.com/in/yzhao062](https://www.linkedin.com/in/yzhao062)  
🏠 <https://viterbi-web.usc.edu/~yzhao010/>  
👤 USC Faculty Directory  
🔍 Google Scholar  
★ 29,000+ GitHub Stars

213-821-2369  
CS Department, GCS Hall  
Los Angeles, CA  
United States, 90089  
Department of Computer Science  
University of Southern California  
Top ~700 Worldwide

## RESEARCH SUMMARY

When something goes wrong, someone has to be able to say what happened and who is answerable. AI systems that act on their own can make that hard: the damage shows up, the path to it often does not. My lab builds three things: tools that catch unsafe permissions before an agent ships, a checkpoint that inspects tool calls before they run, and records good enough to piece together afterward who did what. This work spans **auditable AI systems** across four audit questions: *auditing of action* (auditability and control of agent systems across their lifecycle), *optimization of effort* (agent efficiency, routing, and the cost of defense), *verification of output* (trust and robustness evaluation of foundation models), and *detection in data* (anomaly and out-of-distribution detection). Auditing of action, verification of output, and detection in data ask whether something departed from what was expected and whether anyone can still tell. Optimization of effort asks whether the work was necessary at all. The current frontier is **agent auditability**: deployable open-source infrastructure that audits agents from before deployment to after a run.

### 1. Auditing of Action: Auditability and Control of Agent Systems

When a system acts through tools, APIs, and shared state, what it does becomes the primary failure surface. This thread covers auditability frameworks, runtime control surfaces, agent-specific failure modes, and post-run attribution and recovery, spanning the full lifecycle from pre-deployment to after a run.

- ☐ Agent Auditability
- ☐ Agent Systems
- ☐ Multi-Agent Systems
- ☐ Behavior Verification
- ☐ AI Agent Security
- ☐ Runtime Guardrails
- ☐ Policy Enforcement

### 2. Optimization of Effort: Agent Efficiency, Routing, and the Cost of Defense

The record that makes a system accountable also shows what its effort bought. This thread asks whether the work was necessary: what defense training costs an agent's ability to finish a task, how model choice becomes an online cost-quality decision, and how large-scale detection cuts its compute cost. It is the dimension industry asks about most.

- ☐ Agent Efficiency
- ☐ LLM Routing
- ☐ Cost-Quality Trade-offs
- ☐ Workflow Generation
- ☐ Cost of Defense Training
- ☐ Scalable Detection

### 3. Verification of Output: Trust and Robustness of Foundation Models

Between the input and the action, foundation models add their own failure surface: behavior under attack, hallucination, and adversarial robustness of retrieval. Failures on these surfaces can propagate downstream when model output is passed to an agent or a user-facing system. My lab measures these surfaces through jailbreak detection for vision-language models, causal analysis of hallucination, retrieval-augmented generation attacks, and LLM-as-anomaly-detector benchmarks.

- ☐ Foundation-Model Trust
- ☐ Trustworthiness Benchmarks
- ☐ LLM Hallucination
- ☐ Jailbreak Detection
- ☐ RAG Security
- ☐ Adversarial Robustness
- ☐ Multimodal Trust

#### 4. Detection in Data: Anomaly and Out-of-Distribution Detection

Here the question is whether an input departed from the conditions represented during development and evaluation. This is my methodological and open-source base, built through the PyOD ecosystem and ADBench, and continuing through benchmarks for NLP, automatic OOD detector selection, and modality-specific OOD methods for graphs, multimodal data, and video. When an incident involves distribution shift, these methods supply evidence about whether unfamiliar input conditions contributed to it.

- ☐ Anomaly Detection
- ☐ Outlier Detection
- ☐ Out-of-Distribution Detection
- ☐ Distribution Shift
- ☐ Multimodal OOD
- ☐ Graph OOD
- ☐ Benchmark Design

Tools from this group have been cited in a U.S. Senate committee report, a NIST special publication, and the U.S. Department of Defense Generative AI Toolkit. PyOD was selected by the European Space Agency for spacecraft anomaly detection and is the subject of five published books. TrustLLM serves as an official benchmark in all four editions of the Future of Life Institute AI Safety Index through Summer 2026.

PROFESSIONAL EXPERIENCE    **University of Southern California**  
*Thomas Lord Department of Computer Science*  
Assistant Professor (Tenure-Track) Aug. 2023 - Present

- **F**oundations **O**f **R**obust **T**rustworthy **I**ntelligent **S**ystems (**FORTIS**) Lab: [Link](#)
- USC Machine Learning Center (MaSCle): [Link](#)

#### Fig Work

*AI-powered marketplace for early-career talent*  
Chief Scientific Advisor Jun. 2026 - Present

#### PwC Canada

*Consulting & Deals*  
Senior Consultant (Data Scientist) Aug. 2017 - Jun. 2019  
Consultant (Data Scientist) Feb. 2017 - Jul. 2017  
Research Associate (Intern) May. 2016 - Jan. 2017

EDUCATION    **Carnegie Mellon University** Pittsburgh, PA  
*Ph.D. in Information Systems and Management* Sep. 2019 - May. 2023

- **Affiliation:** CMU automated learning systems group (Catalyst) and Data Analytics Techniques Algorithms (DATA) Lab
- **Advisors and Mentors:** CMU: Prof. Leman Akoglu, Prof. Zhihao Jia, and Prof. George Chen. I collaborate with Prof. Jure Leskovec at Stanford, and Prof. Philip S. Yu at UIC.

**University of Toronto** Toronto, ON  
*Master of Science in Computer Science* Sep. 2015 - Dec. 2016

**University of Cincinnati** Cincinnati, OH  
*Bachelor of Science in Computer Engineering* Sep. 2010 - May. 2015  
**Minor:** *Computer Science and Mathematics*

 AWARDS,  
GRANTS, AND  
FUNDING

### As Principal Investigator (August 2023 onwards)

|                                                               |                           |           |
|---------------------------------------------------------------|---------------------------|-----------|
| NSF FIRE                                                      | <i>Funding</i>            | Aug. 2026 |
| Meta                                                          | <i>Sponsored Research</i> | Aug. 2026 |
| Foresight Institute AI for Safety & Science Nodes             | <i>Gift</i>               | Apr. 2026 |
| NVIDIA Academic Grant Program                                 | <i>Gift</i>               | Mar. 2026 |
| Amazon Research Awards, Fall 2025                             | <i>Gift</i>               | Mar. 2026 |
| USC CCSL Seed Funding                                         | <i>Small Grant</i>        | Jan. 2026 |
| Second Prize CCC Award @ IEEE ICDM, BlueSky Track             | <i>Recognition</i>        | Nov. 2025 |
| Best Short Paper Award @ ACM SIGSPATIAL                       | <i>Recognition</i>        | Nov. 2025 |
| NSF Pathways to Enable Open-Source Ecosystems (POSE), Phase I | <i>Funding</i>            | Aug. 2025 |
| Capital One Research Awards                                   | <i>Grant</i>              | Oct. 2024 |
| Amazon Research Awards, Spring 2024                           | <i>Gift</i>               | Aug. 2024 |
| Best Paper Award @ KDD Resource-Efficient Learning Workshop   | <i>Recognition</i>        | Aug. 2024 |
| NSF ATD                                                       | <i>Funding</i>            | Aug. 2024 |
| NSF POSE Phase II                                             | <i>Funding</i>            | Jun. 2024 |
| Google Cloud Research Innovators                              | <i>Recognition</i>        | Mar. 2024 |
| AAAI New Faculty Highlights                                   | <i>Recognition</i>        | Feb. 2024 |

*Note: Monetary values represent **my portion** of the funding. Total project budgets may be larger.*

### Prior to Principal Investigator Role (Before August 2023)

|                                                    |                    |           |
|----------------------------------------------------|--------------------|-----------|
| Meta 2022 AI4AI Research Award (student co-PI)     | <i>Recognition</i> | Oct. 2022 |
| The Norton Labs Graduate Fellowship                | <i>Fellowship</i>  | Mar. 2022 |
| CMU Presidential Fellowship                        | <i>Fellowship</i>  | 2019      |
| Mitacs-Accelerate Research and Development Funding | <i>Funding</i>     | 2016-2017 |
| University Global Award and Scholarship            | <i>Scholarship</i> | 2010-2015 |
| Mantei/Mae Award & Scholar                         | <i>Award</i>       | 2012-2015 |
| Engineer of the Month                              | <i>Recognition</i> | Jun. 2014 |

*Note: Monetary values are omitted for awards and recognitions received prior to PI role.*

SELECTED  
OPEN-SOURCE  
SOFTWARE,  
BENCHMARKS,  
AND SYSTEMS

Representative open-source research artifacts spanning anomaly detection, trustworthy AI, agent security, and research tooling.

### Primary Contributions

- **PyOD** (*Library*; updated Sep. 2026; **10.0K stars**) PyOD 3 is the most comprehensive Python library for anomaly detection: 60+ detectors across tabular, time series, graph, text, and image data, with an ADEngine orchestration core and an agentic investigation layer driven through natural language. 59M+ downloads. Links: docs — paper.
- **Anomaly-Detection-Resources** (*Resource Hub*; updated Mar. 2026; **9.4K stars**) A curated resource hub for anomaly detection, including papers, benchmarks, tutorials, datasets, tools, and community references.
- **CS-Paper-Checklist** (*Research Toolkit*; updated May. 2026; **1.6K stars**) A practical sanity checklist for improving CS paper quality, structure, and submission readiness.
- **PyGOD** (*Library*; updated Nov. 2024; **1.5K stars**) A Python library for graph outlier detection, with unified APIs and broad coverage of graph anomaly detection algorithms. Links: docs — paper.
- **ADBench** (*Benchmark*; updated Jan. 2026; **1.0K stars**) Official benchmark for anomaly detection with large-scale evaluations across unsupervised, semi-supervised, and supervised settings. Links: paper.
- **agent-style** (*Developer Toolkit*; updated Sep. 2026; **677 stars**) A drop-in rules package that shapes how AI coding and writing agents produce prose at generation time, using 21 curated English style rules (12 from canonical writing authorities; 9 from field observation of LLM output). Works with Claude Code, Codex, Copilot, Cursor, Aider, and other AGENTS.md-compliant tools. Links: docs — demo.
- **combo** (*Library*; updated Sep. 2026; **660 stars**) A Python toolbox for machine learning model and score combination across classification, clustering, and anomaly detection tasks. Links: docs — paper.

- **SUOD** (*Library / System*; updated Sep. 2026; **395 stars**) An acceleration framework for large-scale unsupervised heterogeneous outlier detection with efficient training and inference. Links: [docs](#) — [paper](#).
- **Aegis** (*Security / Compliance Platform*; updated Sep. 2026; **353 stars**) The open-source firewall for AI agents, with pre-execution policy checks, human-in-the-loop approvals, and tamper-evident audit trails for safer and more accountable deployment. Links: [docs](#) — [paper](#) — [demo](#).
- **anywhere-agents** (*Developer Toolkit*; updated Sep. 2026; **244 stars**) One config to rule all your AI agents: portable across every project and session, effective through curated writing, routing, and skills, and safer via a PreToolUse guard that blocks destructive Git and GitHub commands. Supports Claude Code and Codex today, with plans to grow. Links: [docs](#).
- **agent-audit** (*Security Toolkit*; updated Sep. 2026; **229 stars**) An open-source tool for auditing AI agent code and configurations, with OWASP Agentic Top 10 coverage, taint analysis, and MCP configuration auditing. Links: [docs](#) — [paper](#) — [demo](#).
- **awesome-auditable-ai** (*Resource Hub*; updated Sep. 2026; **144 stars**) A curated list of papers, tools, datasets, benchmarks, and standards for building, evaluating, and auditing reliable AI agents, organized around a reader-first reliability map.
- **AD-AGENT** (*Platform*; updated Apr. 2026; **101 stars**) An LLM-driven multi-agent anomaly detection platform for end-to-end workflows across tabular, graph, and time-series settings. Links: [docs](#) — [paper](#).
- **AD-LLM** (*Benchmark*; updated Oct. 2025; **44 stars**) A benchmark for using large language models in anomaly detection, covering zero-shot detection, data augmentation, and model selection. Links: [paper](#).
- **auditable** (*Agent Auditability*; updated Sep. 2026; **23 stars**) An open-source system of record for AI-agent decisions: it captures what each decision relied on, replays it against live state and rolls back the committed action when it no longer holds, and ranks the steps of a finished run by structural blast share, on one typed graph across an agent’s plan, run, and review. Links: [docs](#) — [paper](#) — [demo](#).
- **NLP-ADBench** (*Benchmark*; updated Oct. 2025; **21 stars**) A benchmark for anomaly detection in NLP, including NLPAD datasets and standardized evaluation across end-to-end and embedding-based methods. Links: [paper](#).
- **CatchBench** (*Benchmark*; updated Sep. 2026; **9 stars**) A benchmark for auditing agent failures across the full lifecycle. It scores one auditor’s question against three information states: the declared configuration before a run, a growing prefix of the trace, and the finished trace. The release covers 72 entrants and 1187 declared configurations. Contrasts that do not separate are published unresolved rather than ranked. Links: [paper](#).

## Selected Collaborations

- **TODS** (*System / AutoML*; updated Sep. 2023; **1.7K stars**) An automated time-series outlier detection system with end-to-end pipeline support for preprocessing, feature extraction, model search, and detection. Links: [docs](#) — [paper](#).
- **TDC** (*Benchmark / Platform*; updated Jul. 2025; **1.3K stars**) Therapeutics Data Commons, a platform for AI-ready therapeutic datasets, benchmarks, and evaluation workflows across drug discovery tasks. Links: [docs](#) — [paper](#).
- **TrustLLM** (*Benchmark / Toolkit*; updated Jun. 2025; **633 stars**) A trustworthiness benchmark and toolkit for LLMs; collaboration included follow-on NSF-supported development with the core team. Links: [docs](#) — [paper](#).
- **TrustEval-toolkit** (*Benchmark / Toolkit*; updated Aug. 2025; **137 stars**) A toolkit and benchmark for evaluating trustworthiness of generative foundation models across LLM, VLM, and text-to-image settings. Links: [docs](#) — [paper](#).
- **LangSkills** (*Toolkit / Knowledge Base*; updated Mar. 2026; **116 stars**) An evidence-backed skill library and pipeline for AI agents, with offline searchable bundles distilled from papers and technical sources.

PUBLICATIONS  
 SCHOLAR  
 RESEARCHG

**Preprints & Under Submission** Note: <sup>†</sup>Equal contribution and <sup>♣</sup>Corresponding author for multi-first/corresponding author papers.

### 124. [Yue Zhao](#)

CatchBench: When Can an Agent Failure Be Caught?  
[arXiv preprint arXiv:2608.22808](#)

123. Shawn Li, Wei Yang, Jike Zhong, Jiate Li, Jiawei Yang, You Qin, Ryan Rossi, Franck Dernoncourt, Roger Zimmermann, Yue Wang, Zhengzhong Tu, Vicente Ordonez, Mohit Bansal, [Yue Zhao](#)  
JigShape: Evaluating Visual-Geometric Reasoning in VLMs through Jigsaw Puzzles  
**arXiv preprint arXiv:2607.27670**
122. Yue Huang, Wenjie Wang, Han Bao, Yuchen Ma, Xiaonan Luo, Yi Nian, Haomin Zhuang, Zheyuan Liu, [Yue Zhao](#), Xiangliang Zhang  
MemoHarness: Agent Harnesses That Learn from Experience  
**arXiv preprint arXiv:2607.14159**
121. Tianming Sha, [Yue Zhao](#), Lichao Sun, Yushun Dong  
SkillCenter: A Large-Scale Source-Grounded Skill Library for Autonomous AI Agents  
**arXiv preprint arXiv:2607.07676**
120. Shawn Li, [Yue Zhao](#)  
Agent Safety Is Action Alignment  
**arXiv preprint arXiv:2606.28739**
119. [Yue Zhao](#)  
GRADE: Graph Representation of LLM Agent Dependency and Execution  
**arXiv preprint arXiv:2606.22741**
118. Yingzi Ma, Zhengyue Zhao, Xiaogeng Liu, Minhui Xue, [Yue Zhao](#), Chaowei Xiao  
MaskForge: Structure-Aware Adaptive Attacks for Jailbreaking Diffusion Large Language Models  
**arXiv preprint arXiv:2606.04027**
117. Yuehan Qin, Li Li, Linxin Song, Wei Yang, Jiate Li, Yuqing Yang, [Yue Zhao](#)  
Memory Retrieval for Changing Preferences  
**arXiv preprint arXiv:2606.02976**
116. Ojas Nimase, Jiate Li, [Yue Zhao](#), Yushun Dong  
Can Subgraph Explanations Be Weaponized to Steal Graph Neural Networks?  
**arXiv preprint arXiv:2605.30470**
115. Ojas Nimase, Zhe Chen, Gengpei Qi, [Yue Zhao](#), Xiyang Hu  
GEO-Bench: Benchmarking Ranking Manipulation in Generative Engine Optimization  
**arXiv preprint arXiv:2605.29107**
114. Haolin Chen, Deon Metelski, Leon Qi, Tao Xia, Joonyul Lee, Steve Brown, Kevin Riley, Frank Wang, T. Y. Alvin Liu, Hank Capps MD, Zeyu Tang, Xiangchen Song, Lingjing Kong, Fan Feng, Tianyi Zeng, Zhiwei Liu, Zixian Ma, Hang Jiang, Fangli Geng, Yuan Yuan, Chenyu You, Qingsong Wen, Hua Wei, Yanjie Fu, [Yue Zhao](#), Carl Yang, Biwei Huang, Kun Zhang, Caiming Xiong, Sanmi Koyejo, Eric P. Xing, Philip S. Yu, Weiran Yao  
CHI-Bench: Can AI Agents Automate End-to-End, Long-Horizon, Policy-Rich Healthcare Workflows?  
**arXiv preprint arXiv:2605.16679**
113. Shawn Li, Chenxiao Yu, Han Wang, Wei Yang, Ryan Rossi, Franck Dernoncourt, Xiyang Hu, Philip Yu, Chaowei Xiao, Huan Zhang, [Yue Zhao](#)  
FORTIS: Benchmarking Over-Privilege in Agent Skills  
**arXiv preprint arXiv:2605.09163**
112. Yicheng Gao, Xiaolin Zhou, Yahan Li, [Yue Zhao](#), Ruishan Liu  
MedExAgent: Training LLM Agents to Ask, Examine, and Diagnose in Noisy Clinical Environments  
**arXiv preprint arXiv:2605.07058**
111. Shawn Li, You Qin, Jiate Li, Charith Peris, Lisa Bauer, Roger Zimmermann, [Yue Zhao](#)  
Geometry over Density: Few-Shot Cross-Domain OOD Detection  
**arXiv preprint arXiv:2605.03410**
110. Tiankai Yang, Yi Nian, Xinyuan Li, Ruiyao Xu, Kaize Ding, [Yue Zhao](#)  
Cat-DPO: Category-Adaptive Safety Alignment  
**arXiv preprint arXiv:2604.17299**
109. Tiankai Yang, Jiate Li, Yi Nian, Shen Dong, Ruiyao Xu, Ryan Rossi, Kaize Ding, [Yue Zhao](#)  
No Attacker Needed: Unintentional Cross-User Contamination in Shared-State LLM Agents  
**arXiv preprint arXiv:2604.01350**
108. Haiyue Zhang, Yi Nian, [Yue Zhao](#)  
Agent Audit: A Security Analysis System for LLM Agent Applications  
**arXiv preprint arXiv:2603.22853**

107. Yi Nian, Haosen Cao, Shenzhe Zhu, Henry Peng Zou, Qingqing Luan, Yudi Zhang, Yue Zhao  
When Only the Final Text Survives: Implicit Execution Tracing for Multi-Agent Auditing  
**Under submission**  
**arXiv preprint arXiv:2603.17445**
106. Aojie Yuan, Zhiyuan Su, Yue Zhao  
AEGIS: No Tool Call Left Unchecked – A Pre-Execution Firewall and Audit Layer for AI Agents  
**Under submission**  
**arXiv preprint arXiv:2603.12621**
105. Rujie Ye, Jiayi Zhang, Zhuoxin Liu, Zihao Zhu, Siyuan Yang, Li Li, Tianfu Fu, Franck Dernoncourt, Yue Zhao, Jiacheng Zhu, Ryan Rossi, Wenhao Chai, Zhengzhong Tu  
Agent Banana: High-Fidelity Image Editing with Agentic Thinking and Tooling  
**Under submission**  
**arXiv preprint arXiv:2602.09084**
104. Xiaolin Zhou, Zheng Luo, Yicheng Gao, Qixuan Chen, Xiyang Hu, Yue Zhao, Ruishan Liu  
Fairness or Fluency? An Investigation into Language Bias of Pairwise LLM-as-a-Judge  
**Under submission**  
**arXiv preprint arXiv:2601.13649**
103. Chenxiao Yu, Bowen Yi, Farzan Karimi-Malekabadi, Suhaib Abdurahman, Jinyi Ye, Shrikanth Narayanan, Yue Zhao, Morteza Dehghani  
Tracing Moral Foundations in Large Language Models  
**Under submission**  
**arXiv preprint arXiv:2601.05437**
102. Kay Liu, Yuwei Han, Haoyan Xu, Henry Peng Zou, Yue Zhao, Philip S. Yu  
TAGFN: A Text-Attributed Graph Dataset for Fake News Detection in the Age of LLMs  
**Under submission**  
**arXiv preprint arXiv:2511.21624**
101. Haoyan Xu, Ruizhi Qian, Zhengtao Yao, Ziyi Liu, Li Li, Yuqi Li, Yanshu Li, Wenqing Zheng, Daniele Rosa, Daniel Barcklow, Senthil Kumar, Jieyu Zhao, Yue Zhao  
LLM-Powered Text-Attributed Graph Anomaly Detection via Retrieval-Augmented Reasoning  
**Under submission**  
**arXiv preprint arXiv:2511.17584**
100. Haoyan Xu, Ruizhi Qian, Jiatao Li, Yushun Dong, Minghao Lin, Hanson Yan, Zhengtao Yao, Qinghua Liu, Junhao Dong, Ruopeng Huang, Yue Zhao<sup>✉</sup>, Mengyuan Li<sup>✉</sup>  
A Systematic Study of Model Extraction Attacks on Graph Foundation Models  
**Under submission**  
**arXiv preprint arXiv:2511.11912**
99. Yuexing Hao, Yue Huang, Haoran Zhang, Chenyang Zhao, Zhenwen Liang, Paul Pu Liang, Yue Zhao, Lichao Sun, Saleh Kalantari, Xiangliang Zhang, Marzyeh Ghassemi  
The Role of Computing Resources in Publishing Foundation Model Research  
**Under submission**  
**arXiv preprint arXiv:2510.13621**
98. Wang Wei, Tiankai Yang, Hongjie Chen, Yue Zhao, Franck Dernoncourt, Ryan A. Rossi, Hoda Eldardiry  
Learning to Route LLMs from Bandit Feedback: One Policy, Many Trade-offs  
**Under submission**  
**arXiv preprint arXiv:2510.07429**
97. Langzhou He, Junyou Zhu, Fangxin Wang, Junhua Liu, Haoyan Xu, Yue Zhao, Philip S. Yu, Qitian Wu  
Can Molecular Foundation Models Know What They Don't Know? A Simple Remedy with Preference Optimization  
**Under submission**  
**arXiv preprint arXiv:2509.25509**
96. Yuehan Qin, Li Li, Defu Cao, Tiankai Yang, Yue Zhao  
M3OOD: Automatic Selection of Multimodal OOD Detectors  
**Under submission**  
**arXiv preprint arXiv:2508.11936**
95. Haoyan Xu, Zhengtao Yao, Xuzhi Zhang, Ziyi Wang, Langzhou He, Yushun Dong, Philip S. Yu, Mengyuan Li, Yue Zhao



- GLIP-OOD: Zero-Shot Graph OOD Detection with Foundation Model  
**Under submission**  
**arXiv preprint arXiv:2504.21186**
94. Haoyan Xu, Zhengtao Yao, Ziyi Wang, Zhan Cheng, Xiyang Hu, Mengyuan Li, [Yue Zhao](#)  
 Graph Synthetic Out-of-Distribution Exposure with Large Language Models  
**Under submission**  
**arXiv preprint arXiv:2504.21198**
93. Kaixiang Zhao, Lincan Li, Kaize Ding, Neil Zhenqiang Gong, [Yue Zhao](#), Yushun Dong  
 A Survey of Model Extraction Attacks and Defenses in Distributed Computing Environments  
**Under submission**  
**arXiv preprint arXiv:2502.16065**
92. Shixuan Li, Wei Yang, Peiyu Zhang, Xiongye Xiao, Defu Cao, Yuehan Qin, Xiaole Zhang, [Yue Zhao](#), Paul Bogdan  
 ClimateLLM: Efficient Weather Forecasting via Frequency-Aware Large Language Models  
**Under submission**  
**arXiv preprint arXiv:2502.11059**
91. Lincan Li, Jiaqi Li, Catherine Chen<sup>✉</sup>, Fred Gui<sup>✉</sup>, other collaborators, [Yue Zhao](#)<sup>✉</sup>, Yushun Dong<sup>✉</sup>  
 Political-LLM: Large Language Models in Political Science  
**Under submission**  
**arXiv preprint arXiv:2412.06864**
90. Chenxiao Yu, Jinyi Ye, Yuangang Li, Zhaotian Weng, Zheng Li, Emilio Ferrara, Xiyang Hu<sup>✉</sup>, [Yue Zhao](#)<sup>✉</sup>  
 A Large-Scale Simulation on Large Language Models for Decision-Making in Political Science  
**Under submission**  
**arXiv preprint arXiv:2412.15291**
89. Junda Wu, Hanjia Lyu, Yu Xia, Zhehao Zhang, Joe Barrow, Ishita Kumar, Mehnoosh Mirtahebi, Hongjie Chen, Ryan A. Rossi, Franck Dernoncourt, Tong Yu, Ruiyi Zhang, Jiuxiang Gu, Nesreen K. Ahmed, Yu Wang, Xiang Chen, Hanieh Deilamsalehy, Namyoung Park, Sungchul Kim, Huanrui Yang, Subrata Mitra, Zhengmian Hu, Nedim Lipka, [Yue Zhao](#), Jiebo Luo, Julian McAuley  
 Personalized Multimodal Large Language Models: A Survey  
**Under submission**  
**arXiv preprint arXiv:2412.02142**

### Peer-reviewed Journal Papers

88. Yuehan Qin, Shawn Li, Yi Nian, Xinyan Velocity Yu, [Yue Zhao](#)<sup>✉</sup>, Xuezhe Ma<sup>✉</sup>  
 Don't Let It Hallucinate: Premise Verification via Retrieval-Augmented Logical Reasoning  
*Transactions on Machine Learning Research (TMLR)*, 2026
87. Haoyan Xu, Kay Liu, Zhengtao Yao, Philip S. Yu, Kaize Ding<sup>✉</sup>, [Yue Zhao](#)<sup>✉</sup>  
 LEGO-Learn: Label-Efficient Graph Open-Set Learning  
*Transactions on Machine Learning Research (TMLR)*, 2025
86. Hao Dong, Gaetan Frusque, [Yue Zhao](#), Eleni Chatzi, Olga Fink  
 NNG-Mix: Improving Semi-supervised Anomaly Detection with Pseudo-anomaly Generation  
*IEEE Transactions on Neural Networks and Learning Systems (TNNLS)*, 2024
85. Ling Yang<sup>†</sup>, Zhilong Zhang<sup>†</sup>, Yang Song, Shenda Hong, Runsheng Xu, [Yue Zhao](#), Wentao Zhang, Bin Cui, Ming-Hsuan Yang  
 Diffusion Models: A Comprehensive Survey of Methods and Applications  
*ACM Computing Surveys (CSUR)*, 2023
84. [Yue Zhao](#)<sup>†</sup>, Martin Q. Ma<sup>†</sup>, Xiaorong Zhang, Leman Akoglu  
 The Need for Unsupervised Outlier Model Selection: A Review and Evaluation of Internal Evaluation Strategies  
*ACM SIGKDD Explorations Newsletter (SIGKDD Explor.)*, 2023
83. Kexin Huang<sup>†</sup>, Tianfan Fu<sup>†</sup>, Wenhao Gao<sup>†</sup>, [Yue Zhao](#), Yusuf Roohani, Jure Leskovec, Connor W. Coley, Cao Xiao, Jimeng Sun, Marinka Zitnik  
 Artificial Intelligence Foundation for Therapeutic Science  
*Nature Chemical Biology (NCHEMB)*, 2022

82. Yue Zhao<sup>†</sup>, Zheng Li<sup>†</sup>, Xiyang Hu, Nicola Botta, Cezar Ionescu, George H. Chen  
ECOD: Unsupervised Outlier Detection Using Empirical Cumulative Distribution Functions  
*IEEE Transactions on Knowledge and Data Engineering (TKDE)*, 2022.
81. Yue Zhao, Zain Nasrullah, Zheng Li  
PyOD: A Python Toolbox for Scalable Outlier Detection  
*Journal of Machine Learning Research (JMLR)*, 2019.

## Conference & Workshop Papers

80. Shawn Li, Yue Zhao  
The Autonomy Tax: Defense Training Breaks LLM Agents  
*Conference on Empirical Methods in Natural Language Processing (EMNLP)*, 2026.
79. Wenjie Wang, Yue Huang, Zhengqing Yuan, Han Bao, Shiyi Du, Yuchen Ma, Yue Zhao, Yanfang Ye, Xiangliang Zhang  
SpecAlign: Efficient Specification-Grounded Alignment of Large Language Models via Synthetic Data  
*Conference on Empirical Methods in Natural Language Processing (EMNLP)*, 2026.
78. Yi Nian, Tiankai Yang, Yudi Zhang, Qi Pan, Zelong Xu, Shenzhe Zhu, Qingqing Luan, Yue Huang, Xiangliang Zhang, Yue Zhao  
DOG-DPO: Dynamic Optimization in Geometry for Safety Alignment  
*Findings of the Association for Computational Linguistics: EMNLP*, 2026.
77. Prince Zizhuang Wang, Aojie Yuan, Haiyue Zhang, Xiyang Hu, Yue Zhao, Shuli Jiang  
WeClawArena: An Auditable Sandbox and Benchmark for Cross-User Agents Collaboration and Security in Human-Centered Agent Networks  
*Findings of the Association for Computational Linguistics: EMNLP*, 2026.
76. Zixiang Xu, Yanbo Wang, Yue Huang, Haomin Zhuang, Yujun Zhou, Jiayi Ye, Sixian Li, Zirui Song, Lang Gao, Chenxi Wang, Zhaorun Chen, Wang Pan, Yue Zhao, Jieyu Zhao, Xiangliang Zhang, Xiuying Chen  
Modeling, Evaluating, and Enhancing Reasoning of Large Language Models in Complex Social Environments  
*Findings of the Association for Computational Linguistics: EMNLP*, 2026.
75. Ojas Nimase, Zhengao Li, Xinyu Wei, Jinglin Hu, Luoia Liu, Kangyi Zhao, Yue Zhao, Yushun Dong  
FinanceLLM: A Survey of Large Language Models in Finance  
*Findings of the Association for Computational Linguistics: EMNLP*, 2026.
74. Aojie Yuan, Yi Nian, Haiyue Zhang, Zijian Su, Yue Zhao  
SEVA: Self-Evolving Verification Agent with Process Reward for Fact Attribution  
*ICML 2026 Workshops (AI4GOOD; FAGEN)*.
73. Wang Wei, Hongzheng Yang, Tiankai Yang, Samyadeep Basu, Hongjie Chen, Yue Zhao, Franck Dernoncourt, Ryan A. Rossi, Hoda Eldardiry  
FlexRouter: Learning Complementary Model Sets for Flexible LLM Routing  
*Conference on Language Modeling (COLM)*, 2026.
72. Yiming Tang, Yi Fan, Chenxiao Yu, Tiankai Yang, Yue Zhao, Xiyang Hu  
StealthRank: LLM Ranking Manipulation via Stealthy Prompt Optimization  
*ICML Workshop on Trustworthy AI for Good (AI4GOOD)*, 2026.
71. Han Bao, Yue Huang, Yanbo Wang, Jiayi Ye, Xiangqi Wang, Xiuying Chen, Yue Zhao, Tianyi Zhou, Mohamed Elhoseiny, Xiangliang Zhang  
AutoDavis: Automatic and Dynamic Evaluation Protocol of Large Vision-Language Models on Visual Question-Answering  
*ACM SIGKDD Conference on Knowledge Discovery and Data Mining (KDD Datasets and Benchmarks Track)*, 2026.
70. Zhisheng Qi, Utkarsh Sahu, Li Ma, Haoyu Han, Ryan Rossi, Franck Dernoncourt, Mahantesh Halappanavar, Nesreen Ahmed, Yushun Dong, Yue Zhao, Yu Zhang, Yu Wang  
Benchmarking Knowledge-Extraction Attack and Defense on Retrieval-Augmented Generation  
*ACM SIGKDD Conference on Knowledge Discovery and Data Mining (KDD Datasets and Benchmarks Track)*,  Oral, 2026.



69. Jiate Li, Defu Cao, Li Li, Wei Yang, Yuehan Qin, Chenxiao Yu, Tiannuo Yang, Ryan A. Rossi, Yan Liu, Xiyang Hu, Yue Zhao  
“Someone Hid It”: Query-Agnostic Black-Box Attacks on LLM-Based Retrieval  
*International Conference on Machine Learning (ICML)*, 2026.
68. Shawn Li, Chenxiao Yu, Zhiyu Ni, Hao Li, Charith Peris, Chaowei Xiao, Yue Zhao  
Defenses Against Prompt Attacks Learn Surface Heuristics  
*Annual Meeting of the Association for Computational Linguistics (ACL)*, 2026.
67. Ruiyao Xu, Mihir Parmar, Tiankai Yang, Zhengyu Hu, Yue Zhao, Kaize Ding  
CoAct: Co-Active LLM Preference Learning with Human-AI Synergy  
*Annual Meeting of the Association for Computational Linguistics (ACL)*,  Oral, 2026.
66. Jinbo Liu, Defu Cao, Yifei Wei, Tianyao Su, Yuan Liang, Yushun Dong, Yan Liu, Yue Zhao, Xiyang Hu  
Topology Matters: Measuring Memory Leakage in Multi-Agent LLMs  
*Findings of the Association for Computational Linguistics: ACL*, 2026.
65. Yi Nian, Aojie Yuan, Haiyue Zhang, Jiate Li, Shawn Li, Xiyang Hu, Hua Wei, Xiongye Xiao, Chaowei Xiao, Yue Zhao  
Auditable Agents  
*ACM AI Leadership Summit*, 2026 (condensed 4-page version). Earlier version presented at the *ACL 2026 KnowFM Workshop*.  
**arXiv preprint arXiv:2604.05485**
64. Yixuan Du, Chenxiao Yu, Haoyan Xu, Ziyi Wang, Yue Zhao, Xiyang Hu  
Multimodal Generative Engine Optimization: Exploiting Cross-Modal Knowledge Grounding in VLM-Based Ranking  
*ACL Workshop on Towards Knowledgeable Foundation Models (KnowFM)*, 2026.  
**arXiv preprint arXiv:2601.12263**
63. Shawn Li, Ryan Rossi, Sungchul Kim, Sunav Choudhary, Franck Dernoncourt, Puneet Mathur, Zhengzhong Tu, Yue Zhao  
Charts Are Not Images: On the Challenges of Scientific Chart Editing  
*International Conference on Learning Representations (ICLR)*, 2026
62. Weidi Luo, Tianyu Lu, Qiming Zhang, Xiaogeng Liu, Bin Hu, Yue Zhao, Jieyu Zhao, Song Gao, Patrick McDaniel, Zhen Xiang, Chaowei Xiao  
Doxing via the Lens: Revealing Privacy Leakage in Image Geolocation for Agentic Multi-Modal Large Reasoning Model  
*International Conference on Learning Representations (ICLR)*, 2026
61. Yue Huang, Chujie Gao, Siyuan Wu, Haoran Wang, Xiangqi Wang, Yujun Zhou, Yanbo Wang, Jiayi Ye, Jiawen Shi, Qihui Zhang, Yuan Li, Han Bao, Zhaoyi Liu, Tianrui Guan, Dongping Chen, Ruoxi Chen, other authors, Yue Zhao, other authors, Xiangliang Zhang  
On the Trustworthiness of Generative Foundation Models: Guideline, Assessment, and Perspective  
*International Conference on Learning Representations (ICLR)*, 2026  
<https://trustgen.github.io/>
60. Chengxuan Qian, Shuo Xing, Shawn Li, Yue Zhao, Zhengzhong Tu  
DecAlign: Hierarchical Cross-Modal Alignment for Decoupled Multimodal Representation Learning  
*International Conference on Learning Representations (ICLR)*, 2026
59. Xiong Xiao Xu, Haoran Wang, Yueqing Liang, Philip S. Yu, Yue Zhao, Kai Shu  
Can Multimodal LLMs Perform Time Series Anomaly Detection?  
*The Web Conference (WWW)*, 2026
58. Bo Ni, Yu Wang, Leyao Wang, Branislav Kveton, Franck Dernoncourt, Yu Xia, Hongjie Chen, Reuben Luera, Samyadeep Basu, Subhojyoti Mukherjee, Puneet Mathur, Nesreen K. Ahmed, Junda Wu, Shawn Li, Huixin Zhang, Ruiyi Zhang, Tong Yu, Sungchul Kim, Jiuxiang Gu, Zhengzhong Tu, Alexa Siu, Zichao Wang, Seunghyun Yoon, Nedim Lipka, Namyoung Park, Zihao Lin, Trung Bui, Yue Zhao, Tyler Derr, Ryan A. Rossi  
A Survey on LLM-based Conversational User Simulation  
*Conference of the European Chapter of the Association for Computational Linguistics (EACL)*, 2026
57. Yuangang Li, Yiqing Shen, Yi Nian, Jiechao Gao, Ziyi Wang, Chenxiao Yu, Shawn Li, Jie Wang, Xiyang Hu, Yue Zhao  
Mitigating Hallucinations in Large Language Models via Causal Reasoning  
*Fortieth AAAI Conference on Artificial Intelligence (AAAI)*, 2026

56. Ojas Nimase, Yue Zhao, Yushun Dong  
Navigating Between Explainability and Extractability in Machine Learning as a Service  
*IEEE International Conference on Data Mining (ICDM) BlueSky Track*, 🏆 **Second Prize CCC Award**, 2025.
55. Tiankai Yang, Junjun Liu, Michael Siu, Jiahang Wang, Zhuangzhuang Qian, Chanjuan Song, Cheng Cheng, Xiyang Hu, Yue Zhao  
AD-AGENT: A Multi-agent Framework for End-to-end Anomaly Detection  
*Findings of the Association for Computational Linguistics: IJCNLP-AACL*, 2025.
54. Ruosi Shao, Md Shamim Seraj, Kangyi Zhao, Yingtao Luo, Lincan Li, Bolin Shen, Averi Bates, Yue Zhao, Chongle Pan, Lisa Hightow-Weidman, Shayok Chakraborty, Yushun Dong  
LLM-Empowered Patient-Provider Communication: A Data-Centric Survey From a Clinical Perspective  
*Findings of the Association for Computational Linguistics: IJCNLP-AACL*, 2025.
53. Li Li, Peilin Cai, Ryan A. Rossi, Franck Dernoncourt, Branislav Kveton, Junda Wu, Tong Yu, Lixin Song, Tiankai Yang, Yuehan Qin, Nesreen K. Ahmed, Samyadeep Basu, Subhojyoti Mukherjee, Ruiyi Zhang, Yuxiao Zhou, Zichao Wang, Yue Huang, Yu Wang, Xiangliang Zhang, Philip S. Yu, Xiyang Hu, Yue Zhao  
A Personalized Conversational Benchmark: Towards Simulating Personalized Conversations  
*NeurIPS Workshop on Multi-Turn Interactions in Large Language Models (MTI-LLM)*, 🏆 **Spotlight**, 2025.  
arXiv preprint arXiv:2505.14106
52. Zixiang Xu, Yanbo Wang, Yue Huang, Jiayi Ye, Haomin Zhuang, Zirui Song, Lang Gao, Chenxi Wang, Zhaorun Chen, Yujun Zhou, Sixian Li, Wang Pan, Yue Zhao, Jieyu Zhao, Xiangliang Zhang, Xiuying Chen  
SocialMaze: A Benchmark for Evaluating Social Reasoning in Large Language Models  
*NeurIPS Workshop on Socially Responsible and Trustworthy Foundation Models (ResponsibleFM)*  
arXiv preprint arXiv:2505.23713
51. Yanbo Wang, Zixiang Xu, Yue Huang, Xiangqi Wang, Zirui Song, Lang Gao, Chenxi Wang, Xiangru Tang, Yue Zhao, Arman Cohan, Xiangliang Zhang, Xiuying Chen  
DyFlow: Dynamic Workflow Framework for Agentic Reasoning  
*Advances in Neural Information Processing Systems (NeurIPS)*, 2025
50. Shawn Li, Jiashu Qu, Yuxiao Zhou, Yuehan Qin, Tiankai Yang, Yue Zhao  
Treble Counterfactual VLMs: A Causal Approach to Hallucination  
*Findings of the Association for Computational Linguistics: EMNLP*, 2025.
49. Yuangang Li, Jiaqi Li, Zhuo Xiao, Tiankai Yang, Yi Nian, Xiyang Hu, Yue Zhao  
NLP-ADbench: NLP Anomaly Detection Benchmark  
*Findings of the Association for Computational Linguistics: EMNLP*, 2025.
48. Lincan Li, Eren Erman Ozguven, Yue Zhao, Guang Wang, Yiqun Xie, Yushun Dong  
TyphoFormer: Language-Augmented Transformer for Accurate Typhoon Track Forecasting  
*ACM International Conference on Advances in Geographic Information Systems (SIGSPATIAL)*, 🏆 **Best Short Paper Award**, 2025.
47. Bolin Shen, Eren Erman Ozguven, Yue Zhao, Guang Wang, Yiqun Xie, Yushun Dong  
Learning from the Storm: A Multivariate Machine Learning Approach to Predicting Hurricane-Induced Economic Losses  
*ACM SIGSPATIAL International Workshop on Spatial Intelligence for Smart and Connected Communities (SpatialConnect)*, 2025
46. Yi Nian<sup>†</sup>, Shenzhe Zhu<sup>†</sup>, Yuehan Qin, Shawn Li, Ziyi Wang, Chaowei Xiao, Yue Zhao  
JailDAM: Jailbreak Detection with Adaptive Memory for Vision-Language Model  
*Conference on Language Modeling (COLM)*, 2025.
45. Shawn Li, Peilin Cai, Yuxiao Zhou, Zhiyu Ni, Renjie Liang, You Qin, Yi Nian, Zhengzhong Tu, Xiyang Hu, Yue Zhao  
Secure On-Device Video OOD Detection Without Backpropagation  
*International Conference on Computer Vision (ICCV)*, 2025.
44. Zerui Xu, Fang Wu, Yue Zhao  
Retrieval-Reasoning Large Language Model-based Synthetic Clinical Trial Generation  
*ACM SIGKDD International Conference on Knowledge Discovery and Data Mining (KDD Workshop on AI Agent for Information Retrieval)*, 2025.

- ACM Conference on Bioinformatics, Computational Biology, and Health Informatics (ACM BCB)*, 2025.
43. Haoyan Xu<sup>†</sup>, Zhengtao Yao<sup>†</sup>, Yushun Dong, Ziyi Wang, Ryan A. Rossi, Mengyuan Li, [Yue Zhao](#)  
Few-Shot Graph Out-of-Distribution Detection with LLMs  
*European Conference on Machine Learning and Principles and Practice of Knowledge Discovery in Databases (ECML PKDD)*, 2025.
  42. Tiankai Yang<sup>†</sup>, Yi Nian<sup>†</sup>, Shawn Li, Ruiyao Xu, Yuangang Li, Jiaqi Li, Xiyang Hu, Ryan Rossi, Kaize Ding, Xia Hu, [Yue Zhao](#)  
AD-LLM: Benchmarking Large Language Models for Anomaly Detection  
*Findings of the Association for Computational Linguistics (ACL Findings)*, 2025.
  41. Yu Xia, Subhojyoti Mukherjee, Zhouhang Xie, Junda Wu, Xintong Li, Ryan Aponte, Hanjia Lyu, Joe Barrow, Hongjie Chen, Franck Dernoncourt, Branislav Kveton, Tong Yu, Ruiyi Zhang, Jiuxiang Gu, Nesreen K. Ahmed, Yu Wang, Xiang Chen, Hanieh Deilamsalehy, Sungchul Kim, Zhengmian Hu, [Yue Zhao](#), Nedim Lipka, Seunghyun Yoon, Ting-Hao Kenneth Huang, Zichao Wang, Puneet Mathur, Soumyabrata Pal, Koyel Mukherjee, Zhehao Zhang, Namyong Park, Thien Huu Nguyen, Jiebo Luo, Ryan A. Rossi, Julian McAuley  
From Selection to Generation: A Survey of LLM-based Active Learning  
*Annual Meeting of the Association for Computational Linguistics (ACL)*, 2025.
  40. Kaixiang Zhao, Lincan Li, Kaize Ding, Neil Zhenqiang Gong, [Yue Zhao](#), Yushun Dong  
A Survey on Model Extraction Attacks and Defenses for Large Language Models  
*ACM SIGKDD Conference on Knowledge Discovery and Data Mining (KDD Lecture-Style Tutorial Track)*, 2025.
  39. Shawn Li, Huixian Gong, Hao Dong, Tiankai Yang, Zhengzhong Tu, [Yue Zhao](#)  
DPU: Dynamic Prototype Updating for Multimodal Out-of-Distribution Detection  
*Conference on Computer Vision and Pattern Recognition (CVPR)*, **🏆 Highlight**, 2025
  38. Hanhui Wang, Yihua Zhang, Ruizheng Bai, [Yue Zhao](#), Sijia Liu, Zhengzhong Tu  
Edit Away and My Face Will Not Stay: Personal Biometric Defense against Malicious Generative Editing  
*Conference on Computer Vision and Pattern Recognition (CVPR)*, 2025
  37. Yanbo Wang, Jiayi Ye, Siyuan Wu, Chujie Gao, Yue Huang, Xiuying Chen, [Yue Zhao](#), Xiangliang Zhang  
TRUSTEVAL: A Dynamic Evaluation Toolkit on Trustworthiness of Generative Foundation Models  
*Annual Conference of the North American Chapter of the Association for Computational Linguistics (NAACL Demo Track)*, 2025.
  36. Yuehan Qin<sup>†</sup>, Yichi Zhang<sup>†</sup>, Yi Nian<sup>†</sup>, Xueying Ding, [Yue Zhao](#)  
MetaOOD: Meta-learning for Automatic Out-of-Distribution Detection Model Selection  
*International Conference on Learning Representations (ICLR)*, 2025  
*ACM SIGKDD International Conference on Knowledge Discovery and Data Mining (KDD Workshop on Resource-Efficient Learning for Knowledge Discovery)*, **🏆 Best Paper Award**, 2024.
  35. Sihan Chen, Zhuangzhuang Qian, Michael Siu, Xingcan Hu, Jiaqi Li, Shawn Li, Yuehan Qin, Tiankai Yang, Zhuo Xiao, Wanghao Ye, Yichi Zhang, Yushun Dong, [Yue Zhao](#)  
PyOD 2: A Python Library for Outlier Detection with LLM-powered Model Selection  
*International World Wide Web Conference (The Web Conference Demo Track)*, 2025
  34. Sizhe Liu, Yizhou Lu, Siyu Chen, Xiyang Hu, Jieyu Zhao, Yingzhou Lu, [Yue Zhao](#)  
DrugAgent: Automating AI-aided Drug Discovery Programming through LLM Multi-Agent Collaboration  
*AAAI Workshop on Foundation Models for Biological Discoveries (FMs4Bio)*, 2025.
  33. Hao Dong, [Yue Zhao](#), Eleni Chatzi, Olga Fink  
MultiOOD: Scaling Out-of-Distribution Detection for Multiple Modalities  
*Advances in Neural Information Processing Systems (NeurIPS)*, **🏆 Spotlight**, 2024
  32. Xueying Ding, [Yue Zhao](#), Leman Akoglu  
Fast Unsupervised Deep Outlier Model Selection with Hypernetworks  
*ACM SIGKDD Conference on Knowledge Discovery and Data Mining (KDD)*, 2024
  31. Lichao Sun, Yue Huang, Haoran Wang, Siyuan Wu, Qihui Zhang, Chujie Gao, Yixin Huang, Wenhan Lyu, Yixuan Zhang, Xiner Li, Zhengliang Liu, Yixin Liu, Yijue Wang, Zhikun Zhang, 50+ collaborative authors, [Yue Zhao](#)  
TrustLLM: Trustworthiness in Large Language Models  
*International Conference on Machine Learning (ICML)*, 2024

30. [Yue Zhao](#), Leman Akoglu  
Hyperparameter Optimization for Unsupervised Outlier Detection  
*International Conference on Automated Machine Learning (AutoML)*, 2024
29. [Yue Zhao](#)  
Towards Reproducible, Automated, and Scalable Anomaly Detection  
*AAAI Conference on Artificial Intelligence (AAAI), New Faculty Highlights*, 2024
28. Minqi Jiang<sup>†</sup>, Chaochuan Hou<sup>†</sup>, Ao Zheng<sup>†</sup>, Songqiao Han, Hailiang Huang<sup>♣</sup>, Qingsong Wen, Xiyang Hu<sup>♣</sup>, [Yue Zhao](#)<sup>♣</sup>  
ADGym: Design Choices for Deep Anomaly Detection.  
*Advances in Neural Information Processing Systems (NeurIPS)*, 2023  
(<sup>♣</sup>Corresponding author)
27. Jaemin Yoo, [Yue Zhao](#), Lingxiao Zhao, Leman Akoglu  
DSV: An Alignment Validation Loss for Self-supervised Outlier Model Selection  
*European Conference on Machine Learning and Principles and Practice of Knowledge Discovery in Databases (ECML/PKDD)*, 2023
26. Peng Xu, Lin Zhang, Xuanzhou Liu, Jiaqi Sun, [Yue Zhao](#), Haiqin Yang, Bei Yu  
Do Not Train It: A Linear Neural Architecture Search of Graph Neural Networks  
*International Conference on Machine Learning (ICML)*, 2023
25. [Yue Zhao](#), Guoqing Zheng, Subhabrata Mukherjee, Robert McCann, Ahmed Awadallah  
ADMoe: Anomaly Detection with Mixture-of-Experts from Noisy Labels  
*Thirty-Seventh AAAI Conference on Artificial Intelligence (AAAI)*, 2023
24. [Yue Zhao](#), George H. Chen, Zhihao Jia  
TOD: GPU-accelerated Outlier Detection via Tensor Operations  
*International Conference on Very Large Data Bases (VLDB)*, 2023
23. Songqiao Han<sup>†</sup>, Xiyang Hu<sup>†</sup>, Hailiang Huang<sup>†</sup>, Minqi Jiang<sup>†</sup>, [Yue Zhao](#)<sup>†♣</sup>  
ADBench: Anomaly Detection Benchmark  
*Advances in Neural Information Processing Systems (NeurIPS)*, 2022  
(<sup>†</sup>Equal contribution & <sup>♣</sup>Corresponding author)
22. [Yue Zhao](#)<sup>†</sup>, Kay Liu<sup>†</sup>, Yingdong Dou<sup>†</sup>, et al.  
Benchmarking Node Outlier Detection on Graphs  
*Advances in Neural Information Processing Systems (NeurIPS)*, 2022.
21. [Yue Zhao](#), Xiaorong Zhang, Leman Akoglu  
ELECT: Toward Unsupervised Outlier Model Selection  
*IEEE International Conference on Data Mining (ICDM)*, 2022.
20. Zhiming Xu, Xiao Huang, [Yue Zhao](#), Yushun Dong, Jundong Li  
Contrastive Attributed Network Anomaly Detection with Data Augmentation  
*Pacific-Asia Conference on Knowledge Discovery and Data Mining (PAKDD)*, 2022.
19. [Yue Zhao](#), Ryan A. Rossi, Leman Akoglu  
Automatic Unsupervised Outlier Model Selection  
*Advances in Neural Information Processing Systems (NeurIPS)*, 2021.
18. Kwei-Hereng Lai, Daochen Zha, Junjie Xu, [Yue Zhao](#), Guanchu Wang, Xia Hu  
Revisiting Time Series Outlier Detection: Definitions and Benchmarks  
*Advances in Neural Information Processing Systems (NeurIPS)*, 2021
17. Kexin Huang<sup>†</sup>, Tianfan Fu<sup>†</sup>, Wenhao Gao<sup>†</sup>, [Yue Zhao](#), Yusuf Roohani, Jure Leskovec, Connor W. Coley, Cao Xiao, Jimeng Sun, Marinka Zitnik  
Therapeutics Data Commons: Machine Learning Datasets and Tasks for Drug Discovery and Development  
*Advances in Neural Information Processing Systems (NeurIPS)*, 2021
16. [Yue Zhao](#)<sup>†</sup>, Xiyang Hu<sup>†</sup>, Cheng Cheng, Cong Wang, Changlin Wan, Wen Wang, Jianing Yang, Haoping Bai, Zheng Li, Cao Xiao, Yunlong Wang, Zhi Qiao, Jimeng Sun, Leman Akoglu  
SUOD: Accelerating Large-scale Unsupervised Heterogeneous Outlier Detection  
*Conference on Machine and Learning Systems (MLSys)*, 2021.
15. Kwei-Hereng Lai<sup>†</sup>, Daochen Zha<sup>†</sup>, Guanchu Wang, Junjie Xu, [Yue Zhao](#), Devesh Kumar, Yile Chen, Purav Zumkhawaka, Mingyang Wan, Diego Martinez and Xia Ben Hu  
TODS: An Automated Time Series Outlier Detection System (Demo paper)  
*Thirty-Fifth AAAI Conference on Artificial Intelligence (AAAI)*, 2021.

14. Meng-Chieh Lee, Yue Zhao, Aluna Wang, Pierre Jinghong Liang, Leman Akoglu, Vincent S. Tseng, Christos Faloutsos  
AutoAudit: Mining Accounting and Time-Evolving Graphs  
*IEEE International Conference on Big Data (Big Data)*, 2020
13. Changlin Wan, Dongya Jia, Yue Zhao, Wennan Chang, Sha Cao, Xiao Wang, and Chi Zhang  
A Data Denoising Approach to Optimize Functional Clustering of Single Cell RNA-sequencing Data  
*IEEE International Conference on Bioinformatics and Biomedicine (BIBM)*, 2020
12. Yue Zhao, Xueying Ding, Jianing Yang, Haoping Bai.  
SUOD: Toward Scalable Unsupervised Outlier Detection  
**Workshops at the Thirty-Fourth AAAI Conference on Artificial Intelligence**, 2020.  
**Extended version published in *MLSys 2021*.**
11. Zheng Li, Yue Zhao, Nicola Botta, Cezar Ionescu, Xiyang Hu  
COPOD: Copula-Based Outlier Detection  
*IEEE International Conference on Data Mining (ICDM)*, 2020.
10. Zheng Li, Yue Zhao, Jialin Fu  
SYNC: A Copula based Framework for Generating Synthetic Data from Aggregated Sources  
*IEEE International Conference on Data Mining Workshops (ICDMW)*, 2020.
9. Yiqun Mei, Yue Zhao, Wei Liang  
DSR: An Accurate Single Image Super Resolution Approach for Various Degradations  
*IEEE International Conference on Multimedia and Expo (ICME)*, 2020, London, UK.
8. Yue Zhao, Xuejian Wang<sup>†</sup>, Cheng Cheng<sup>†</sup>, Xueying Ding<sup>†</sup>  
Combining Machine Learning Models and Scores using combo Library (Demo paper)  
*Thirty-Fourth AAAI Conference on Artificial Intelligence (AAAI)*, 2020.
7. Colin Wan, Zheng Li, Alicia Guo, Yue Zhao  
SynC: A Unified Framework for Generating Synthetic Population with Gaussian Copula  
**Workshops at the Thirty-Fourth AAAI Conference on Artificial Intelligence**, 2020.  
**Extended version published in *ICDMW 2020*.**
6. Zain Nasrullah, Yue Zhao  
Music Artist Classification with Convolutional Recurrent Neural Networks  
*IEEE International Joint Conference on Neural Networks (IJCNN)*, 2019, Hungary.
5. Yue Zhao, Zain Nasrullah, Maciej K. Hryniewicki, Zheng Li  
LSCP: Locally Selective Combination in Parallel Outlier Ensembles  
*SIAM International Conference on Data Mining (SDM)*, 2019, Calgary, Canada.
4. Yue Zhao, Maciej K. Hryniewicki  
DCSO: Dynamic Combination of Detector Scores for Outlier Ensembles  
*ACM SIGKDD International Conference on Knowledge Discovery and Data Mining (KDD Workshop on Outlier Detection De-constructed)*, 2018, London, UK.  
**Extended version published in *SDM 2019*, renamed to LSCP.**
3. Yue Zhao, Maciej K. Hryniewicki  
XGBOD: Improving Supervised Outlier Detection with Unsupervised Representation Learning  
*IEEE International Joint Conference on Neural Networks (IJCNN)*, 2018, Rio, Brazil.
2. Yue Zhao, Maciej K. Hryniewicki, Francesca Cheng, Boyang Fu, Xiaoyu Zhu  
Employee Turnover Prediction with Machine Learning: A Reliable Approach  
*Intelligent System Conference (Intellisys)*, 2018, London, UK.
1. Yue Zhao<sup>†</sup>, Zhongtian Qiu<sup>†</sup>, Yiqing Yang<sup>†</sup>, Weiwei Li<sup>†</sup>, Mingming Fan  
An Empirical Study of Touch-based Authentication Methods on Smartwatches  
*ACM International Symposium on Wearable Computers (ISWC)*, 2017, Maui, USA.

INTERNSHIP  
EXPERIENCE

**NortonLifeLock Research Group**  
Machine Learning Research Intern

2022

**Microsoft Research**  
Machine Learning Research Intern

2022

**Stanford University, Computer Science Department**  
Visiting Student Researcher (Prof. Jure Leskovec)

2021

**IQVIA, Analytics Center of Excellence**  
Machine Learning Research Intern

2020

**Siemens PLM Software USA**  
Software Engineer (Intern & Contract)

Mar. 2012 - Dec. 2014

TEACHING  
EXPERIENCE

**University of Southern California**

Los Angeles, CA

**Instructor**

Fall 2026 (scheduled)

*CSCI 566 Deep Learning and Its Applications*

**Instructor**

Spring 2026

*CSCI 699 Adversarial and Trustworthy Foundation Models*

**Instructor**

Spring 2025

*CSCI 566 Deep Learning and Its Applications*

**Instructor**

Spring 2024

*CSCI 566 Deep Learning and Its Applications*

**Carnegie Mellon University**

Pittsburgh, PA

**Teaching Assistant**

Fall 2022

*Managing Digital Business* (Prof. David Riel)

**Teaching Assistant & co-Instructor** (lectures on AutoML and MLSys)

Spring 2022 – Fall 2020

*Intro to Artificial Intelligence* (Prof. David Steier)

**Teaching Assistant**

Spring 2022

*Digital Transformation* (Prof. David Riel)

**Teaching Assistant** (helping on course topics)

Fall 2021

*Statistics for IT Managers* (Prof. Daniel Nagin)

**University of Toronto**

Toronto, ON

**Teaching Assistant & Lab Session Instructor**

Fall 2015

*Embedded Systems* (Prof. Philip Anderson)

**University of Cincinnati**

Cincinnati, OH

**Teaching Assistant & Lab Session Instructor**

Fall 2014

*Intro to Programming* (Prof. George Purdy)

PH.D. STUDENTS

- Haoyan Xu (USC, ECE Ph.D. Candidate, 2024 Spring-), co-advised by Mengyuan Li,  **Capital One Fellowship**
- Yuehan Qin (USC, CS Ph.D. Candidate, 2024 Fall-)
- Tiankai Yang (USC, CS Ph.D., 2024 Fall-)
- Shawn Li (USC, CS Ph.D. Candidate, 2024 Fall-),  **Capital One Fellowship**, **Amazon ML Fellowship**
- Jiate Li (USC, CS Ph.D., 2025 Fall-)
- Wei Yang (USC, CS Ph.D. Candidate, 2026 Fall-), co-advised by Jesse Thomason and Xuezhe Ma
- Yi Nian (USC, CS Ph.D., incoming)
- Chenxiao Yu (USC, CS Ph.D., 2026 Fall-)
- Zixiang Xu (USC, CS Ph.D., 2026 Fall-)

STUDENT  
COMMITTEE



| Term        | Student                                | Committee              |
|-------------|----------------------------------------|------------------------|
| Fall 2026   | Chen Chu (USC, CS Ph.D.)               | Qualification          |
|             | Yuqing Yang (USC, CS Ph.D.)            | Qualification          |
|             | Zeyu Shangguan (USC, CS Ph.D.)         | Qualification          |
|             | Chufan Shi (USC, CS Ph.D.)             | Qualification          |
|             | Linxin Song (USC, CS Ph.D.)            | Qualification          |
|             | Raja Kumar (USC, CS Ph.D.)             | Qualification          |
|             | Ziyi Liu (USC, CS Ph.D.)               | Thesis                 |
| Summer 2026 | Defu Cao (USC, CS Ph.D.)               | Defense                |
| Spring 2026 | Longchao Da (ASU, CS Ph.D.)            | Defense                |
|             | Yahan Li (USC, CS Ph.D.)               | Qualification          |
|             | Jiageng Mao (USC, CS Ph.D.)            | Thesis                 |
|             | Xinyue Cui (USC, CS Ph.D.)             | Qualification          |
|             | Jiawei Yang (USC, CS Ph.D.)            | Qualification + Thesis |
|             | Elnaz Rahmati (USC, CS Ph.D.)          | Qualification          |
| Fall 2025   | Hao Dong (ETH Zurich, CS Ph.D.)        | Defense                |
|             | Haonan Wang (USC, ECE Ph.D.)           | Defense                |
| Spring 2025 | Mengxi Wu (USC, CS Ph.D.)              | Qualification          |
|             | Xingrui Wang (USC, ME Ph.D.)           | Defense                |
| Fall 2024   | Alex Bisberg (USC, CS Ph.D.)           | Thesis                 |
| Summer 2024 | Gengyu Rao (USC, CS Ph.D.)             | Defense                |
|             | Mehrdad Kiamari (USC, ECE Ph.D.)       | Defense                |
| Spring 2024 | Maria Despoina Siampou (USC, CS Ph.D.) | Qualification          |
|             | Yuan Meng (USC, ECE Ph.D.)             | Defense                |
|             | Hassan Hamad (USC, ECE Ph.D.)          | Qualification          |
|             | Yizhou Zhang (USC, CS Ph.D.)           | Thesis                 |
|             | Haoming Li (USC, CS Ph.D.)             | Qualification          |
|             | Arash Hajisafi (USC, CS Ph.D.)         | Qualification          |
| Fall 2023   | Yi Chien Lin (USC, ECE Ph.D.)          | Qualification          |
|             | Yuke Zhang (USC, ECE Ph.D.)            | Qualification          |

## SERVICES

### Institutional Leadership

- Associate Co-Director, USC Institute on Ethics and Trust in Computing (IETC), 2026–2027

### Community Outreach

- Founder & Maintainer, CSpHD.org: a non-profit community for CS/AI/EE/Stats PhDs (since 2022, several thousand members, 500+ PhD/RA/postdoc openings), supporting students from before enrollment to the faculty job market.

### Conference/Workshop Organizing Roles

- Doctoral Consortium Co-Chair, AAAI 2027
- Future Faculty Highlights Co-Chair, IEEE BigData 2026
- Publicity Co-Chair, WSDM 2027
- Workflow Co-Chair, KDD 2023
- Co-organizer, SURGeLLM: Structured Understanding, Retrieval, and Generation in LLMs era Workshop @ ACL 2026
- Co-organizer, AgentSkills '26: The First Workshop on Agent Skills @ CAIS Conference 2026
- Co-organizer, AI for Financial Fraud Detection & Prevention Workshop @ 6th ACM International Conference on AI in Finance

### Funding Proposal Review Roles

- National Science Foundation (NSF)
- Dutch Research Council (NWO)

### **Journal Editorial Roles**

- Action Editor, Transactions on Machine Learning Research (TMLR), 2026–present
- Associate Editor, ACM Transactions on AI for Science (TAIS), 2025–present
- Associate Editor, IEEE Transactions on Neural Networks and Learning Systems (TNNLS), 2024–present
- Action Editor, Journal of Data-centric Machine Learning Research (DMLR), 2024–present

### **Conference/Workshop AC and Reviewer Roles**

- ICLR 2025 (AC), ICLR 2026 (AC)
- AAAI 2021, 2022, 2023, 2025 (Senior PC), 2026 (Senior PC)
- ICML 2024, 2025 (AC), 2026 (AC)
- NeurIPS 2021, 2022, 2023, 2025 (AC), 2026 (AC)
- AISTATS 2024, 2025 (AC)
- MLSys 2024, 2026
- KDD 2020, 2021, 2022, 2023, 2026
- ACL Rolling Review (ARR) 2026
- IJCAI 2022, 2023
- AAAI Demonstrations 2021, 2022
- MICCAI 2020, 2021, 2022
- ICDM 2020
- KDD Workshop on Outlier Detection and Description (ODD), 2021
- KDD Workshop on Anomaly and Novelty Detection (ANDEA), 2021, 2022
- IJCAI Workshop on Artificial Intelligence for Anomalies and Novelties (AI4AN), 2020, 2021
- INFORMS Workshop on Data Science 2021

### **Journal Review Roles**

- Journal of Machine Learning Research (JMLR)
- PNAS Nexus
- Machine Learning
- IEEE Transactions on Pattern Analysis and Machine Intelligence (TPAMI)
- IEEE Transactions on Knowledge and Data Engineering (TKDE)
- IEEE Internet of Things Journal (IoT-J)
- IEEE Intelligent Systems
- IEEE Journal on Selected Areas in Communications (J-SAC)
- Data Mining and Knowledge Discovery (DMAI)
- ACM Transactions on Management Information Systems (TMIS)
- Knowledge and Information Systems (KAIS)
- INFORMS Journal on Computing (IJOC)
- Big Data
- Artificial Intelligence Review (AIRE)
- Neurocomputing
- IEEE Transactions on Systems, Man, and Cybernetics: Systems
- IEEE/ACM Transactions on Computational Biology and Bioinformatics (TCBB)
- IEEE Network Magazine
- IEEE Computational Intelligence Magazine (CIM)
- BioData Mining
- European Journal of Management and Business Economics (EJMBE)
- The Journal of Open Source Software (JOSS)

|                    |                                             |                                                                                      |           |
|--------------------|---------------------------------------------|--------------------------------------------------------------------------------------|-----------|
| TALKS AND LECTURES | NVIDIA Agent & Security                     | <i>Safe Models, Broken Agents: How Deployed LLM Agents Break Without an Attacker</i> | May. 2026 |
|                    | Amazon AI Security                          | <i>Safe Models, Broken Agents: How Deployed LLM Agents Break Without an Attacker</i> | May. 2026 |
|                    | USC symposium on Frontiers of ML/AI         | <i>Towards Robust AI: Advances in Outlier and OOD Detection</i>                      | Mar. 2025 |
|                    | NUS Tea Talk                                | <i>Towards Robust AI: Advances in Outlier and OOD Detection</i>                      | Jan. 2025 |
|                    | SFU@NeurIPS'24                              | <i>Towards Robust AI: Advances in Outlier and OOD Detection</i>                      | Dec. 2024 |
|                    | KAIST                                       | <i>Unsupervised Model Selection: Automation with Meta-learning and LLMs</i>          | Nov. 2024 |
|                    | Kennesaw State University                   | <i>Unsupervised Model Selection: Automation with Meta-learning and LLMs</i>          | Oct. 2024 |
|                    | LinkedIn Anti-Abuse AI                      | <i>Outlier Detection: Automation, Systems, and GenAI</i>                             | Aug. 2024 |
|                    | Amazon Security AI                          | <i>Outlier Detection: Automation, Systems, and GenAI</i>                             | Aug. 2024 |
|                    | New York University                         | <i>Outlier Detection: Automation, Systems, and GenAI</i>                             | Aug. 2024 |
|                    | University of Washington                    | <i>Outlier Detection: Automation, Systems, and GenAI</i>                             | Jun. 2024 |
|                    | Microsoft                                   | <i>Outlier Detection: Automation, Systems, and GenAI</i>                             | Jun. 2024 |
|                    | USC Retreat on AI and Engineering Safety    | <i>Safety Measures for LLMs</i>                                                      | Apr. 2024 |
|                    | Visa Research                               | <i>Towards Reproducible, Automated, and Scalable AD</i>                              | Apr. 2024 |
|                    | USC Symposium on Frontiers of Generative AI | <i>Generative AI for Anomaly Detection</i>                                           | Mar. 2024 |
|                    | AAAI New Faculty Highlights (invited)       | <i>Towards Reproducible, Automated, and Scalable AD</i>                              | Feb. 2024 |
|                    | U of Nevada, Las Vegas                      | <i>Automated and Scalable ML Algorithms and Systems</i>                              | Oct. 2023 |
|                    | Samsung Seminar                             | <i>Automated and Scalable Anomaly Detection Systems</i>                              | Aug. 2023 |
|                    | KDD SoCal Day                               | <i>Enable Applications by ML with Noisy Inputs</i>                                   | Aug. 2023 |
|                    | CMU Catalyst                                | <i>How (Not) to Fail Your Academic Job Search</i>                                    | May. 2023 |
|                    | KAUST                                       | <i>Automated and Scalable ML Algorithms and Systems</i>                              | Apr. 2023 |
|                    | Emory University                            | <i>Automated and Scalable ML Algorithms and Systems</i>                              | Apr. 2023 |
|                    | USC                                         | <i>Automated and Scalable ML Algorithms and Systems</i>                              | Mar. 2023 |
|                    | UC Davis                                    | <i>Automated and Scalable ML Algorithms and Systems</i>                              | Mar. 2023 |
|                    | Stony Brook University                      | <i>Automated and Scalable ML Algorithms and Systems</i>                              | Feb. 2023 |
|                    | University of Chicago                       | <i>Automated and Scalable ML Algorithms and Systems</i>                              | Feb. 2023 |
|                    | UC Merced                                   | <i>Automated and Scalable ML Algorithms and Systems</i>                              | Feb. 2023 |
|                    | CMU PDL Meeting                             | <i>Automated and Scalable ML Algorithms and Systems</i>                              | Jan. 2023 |
|                    | CMU Data Science Seminar                    | <b>Guest Lecture</b> <i>Automated Anomaly Detection</i>                              | Nov. 2022 |
|                    | LoG Seminar                                 | <i>Large-scale Graph Anomaly Detection</i>                                           | Oct. 2022 |
|                    | Intuit                                      | <i>Anomaly Detection for Financial Risk Modeling</i>                                 | Aug. 2022 |
|                    | Rice University                             | <i>Large-scale Anomaly Detection with Automation</i>                                 | Sep. 2022 |
|                    | Microsoft Research                          | <i>Weakly-supervised Anomaly Detection</i>                                           | Sep. 2022 |
|                    | Wells Fargo                                 | <i>Anomaly Detection for Financial Risk Modeling</i>                                 | Aug. 2022 |
|                    | Columbia University                         | <b>Guest Lecture</b> <i>Anomaly Detection</i>                                        | Jul. 2022 |
|                    | Morgan Stanley                              | <i>Automated Outlier Detection</i>                                                   | Jun. 2022 |
|                    | Microsoft Research                          | <i>Automated Outlier Detection</i>                                                   | Jun. 2022 |
|                    | Morgan Stanley                              | <i>Large-scale Anomaly Detection Systems</i>                                         | Mar. 2022 |
|                    | Rutgers Business School                     | <i>Outlier Model Selection</i>                                                       | Mar. 2022 |
|                    | Tesla                                       | <i>Large-scale Anomaly Detection Systems</i>                                         | Feb. 2022 |
|                    | Catalyst, CMU                               | <i>Systems for Data Mining Algorithms</i>                                            | Dec. 2021 |
|                    | E&Y Canada                                  | <i>ML applications in Data Analytics</i>                                             | Oct. 2021 |
|                    | University of Nottingham                    | <i>General Machine Learning Applications</i>                                         | Jan. 2021 |